

PROJECT 3 INSTRUCTIONS

Show and Tell

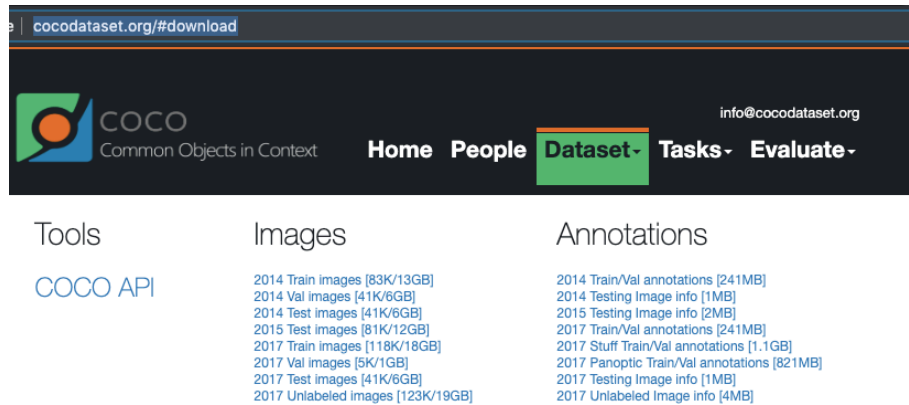
1. Environment Setup
 - a. Install essential libraries
Python3 and upper, Pytorch, torchvision, nltk, numpy, tensorboardX, scikit-image, CUDA

Go to Pytorch official website and install the library proper for your environment:
<https://pytorch.org/>
The same thing with CUDA, be sure to pair with your NAVIDA driver version.
 - b. Use virtual environment if necessary, simply one should be mini conda, Docker and Singularity may be used as well. For first time user, go on the websites to download pre-built img or container with essentials libraries you need. For professionals, you may build your own.
2. Dataset and Json files Preparation (Command Line or interactive)
 - a. git clone <https://github.com/AaronCCWong/Show-Attend-and-Tell.git>
 - b. Download the pre-trained models in the website above, and there are four available there. Put them in the model folder.

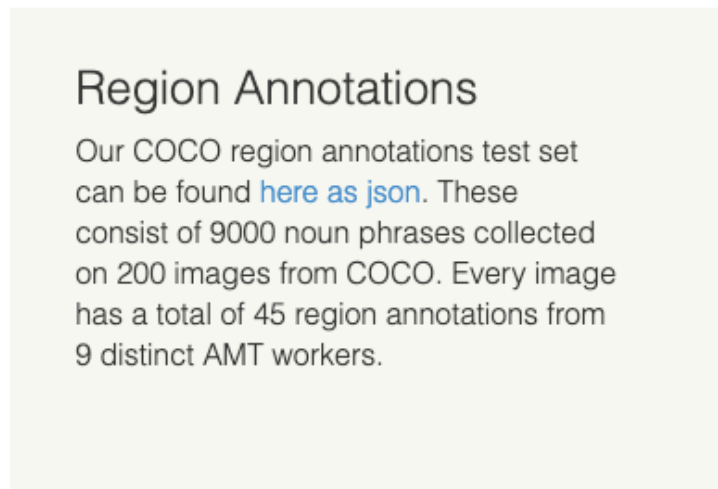
A PyTorch implementation

For a trained model to load into the decoder, use

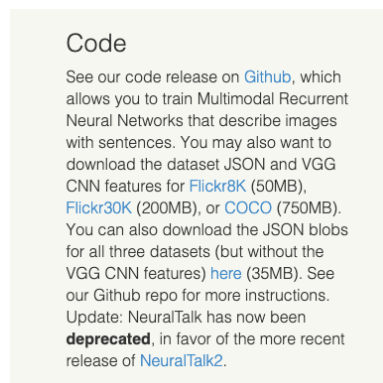
- [VGG19](#)
 - [ResNet152](#)
 - [ResNet152 No Teacher Forcing](#)
 - [VGG19 No Gating Scalar](#)
- c. Download data from the <http://cocodataset.org/#download>, and put them in the corresponding folders as mentioned on github.



- d. Download COCO dataset split JSON file from the link <https://cs.stanford.edu/people/karpathy/deepimagesent/>, and rename it as dataset.json in the corresponding folders as mentioned on github.



- e. Run preprocessing to create JSON files by command line: Python generate_json_data.py
- f. Download other JSON files if you would like to try training the dataset by yourself. (Don't have to do this if you just want to use the pre-trained to have a prediction.)



3. Run the prediction
Follow the instruction on the github or you may refer to my command line for reference as well.

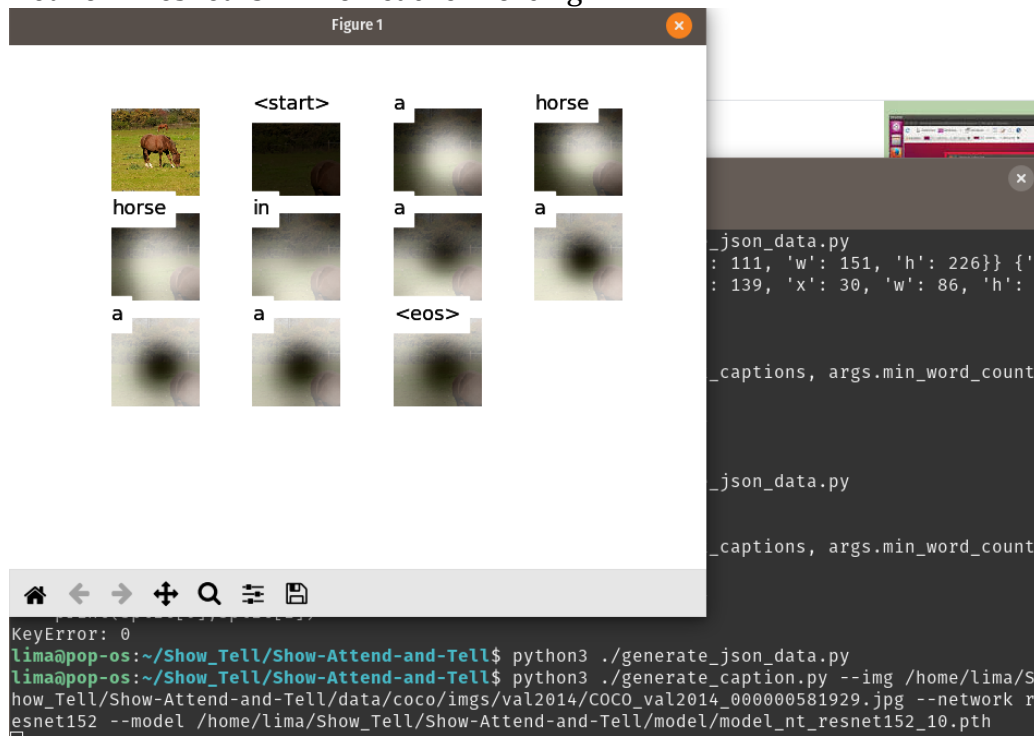
Reference Results:

The implement was performed on a Linux Ubuntu 18.04 LTS system.

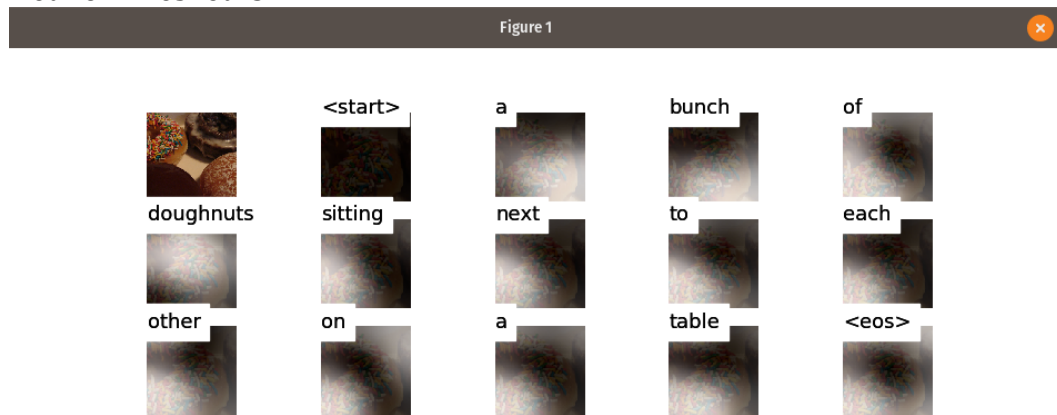
```
lima@pop-os:~/Show_Tell$ python3 --version
Python 3.6.8
lima@pop-os:~/Show_Tell$ python3 ./lib_verification.py
torch: 1.3.0
torchvision 0.4.1
numpy 1.17.2
tensorboardX 1.9
nltk 3.4.5
```

scikit-image 3.4.5

Network: resnet152 – No Teacher Forcing



Network: resnet152



```
lima@pop-os:~/Show_Tell/Show-Attend-and-Tell$ python3 ./generate_caption.py --img /home/lima/Show_Tell/Show-Attend-and-Tell/data/coco/imgs/val2014/COCO_val2014_000000581913.jpg --network resnet152 --model /home/lima/Show_Tell/Show-Attend-and-Tell/model/model_resnet152_10.pth
```

Network: vgg19



TensorFlow Implement:

https://www.tensorflow.org/tutorials/text/image_captioning

Click Run in Google Colab directly; don't copy it into your own colab, which will limit the capacity of the RAM and Disk.

Training Procedure:


```
[30] Epoch 17 Batch 200 Loss 0.2982
Epoch 17 Batch 300 Loss 0.3104
Epoch 17 Loss 0.307570
Time taken for 1 epoch 123.64052319526672 sec
```

```
Epoch 18 Batch 0 Loss 0.3314
Epoch 18 Batch 100 Loss 0.3409
Epoch 18 Batch 200 Loss 0.2828
Epoch 18 Batch 300 Loss 0.2877
Epoch 18 Loss 0.288281
Time taken for 1 epoch 118.91653037071228 sec
```

```
Epoch 19 Batch 0 Loss 0.2945
Epoch 19 Batch 100 Loss 0.2900
Epoch 19 Batch 200 Loss 0.2545
Epoch 19 Batch 300 Loss 0.2624
Epoch 19 Loss 0.268873
Time taken for 1 epoch 132.19374108314514 sec
```

```
Epoch 20 Batch 0 Loss 0.2902
Epoch 20 Batch 100 Loss 0.2276
Epoch 20 Batch 200 Loss 0.2581
Epoch 20 Batch 300 Loss 0.2787
Epoch 20 Loss 0.252668
Time taken for 1 epoch 120.01273250579834 sec
```

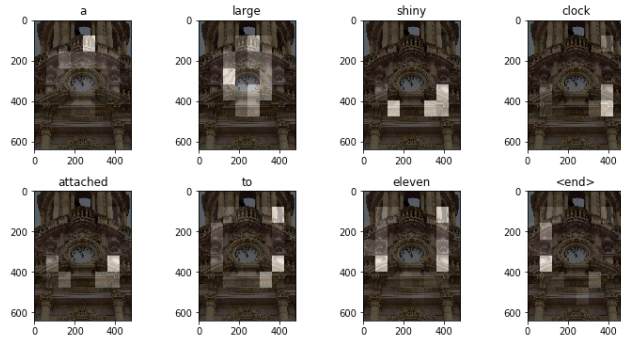
```
[31] plt.plot(loss_plot)
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Loss Plot')
plt.show()
```



Caption Results:

[34]

Real Caption: <start> a clock that is in a very nice building <end>
 Prediction Caption: a large shiny clock attached to eleven <end>



```
image_url = 'https://tensorflow.org/images/surf.jpg'
image_extension = image_url[-4:]
image_path = tf.keras.utils.get_file('image'+image_extension,
                                     origin=image_url)
```

```
result, attention_plot = evaluate(image_path)
print ('Prediction Caption:', ' '.join(result))
plot_attention(image_path, result, attention_plot)
# opening the image
Image.open(image_path)
```

Downloading data from <https://tensorflow.org/images/surf.jpg>
 65536/64400 [=====] - 0s 0us/step
 Prediction Caption: a man riding a surfboard riding a surfboard <end>

